

Week 8: Hypothesis Testing, Correlation vs. Causation

DSUA111: Data Science for Everyone, NYU, Fall 2020

TA Jeff, jpj251@nyu.edu

- This slideshow: <https://jjacobs.me/dsua111-sections/week-08>
(<https://jjacobs.me/dsua111-sections/week-08>).
- All materials: <https://github.com/jpowerj/dsua111-sections>
(<https://github.com/jpowerj/dsua111-sections>).

Outline

1. Hypothesis Testing Overview
2. Testing Coins
3. Null vs. Alternative Hypotheses
4. Test Statistics
5. The Normal Distribution
6. Correlation vs. Causation

Hypothesis Testing Overview

tl;dr

- If your theory was true, what would the data look like?
- Now compare that to the actual, **observed** data

Testing Coins

Example: I walk up to you and say "Hey, wanna gamble? We'll each put in a dollar, then I'll flip this coin. Heads I get the \$2, tails you get the \ \$2"

- Xavier's Theory: I think the coin is fair. Heads and tails will come up about the same number of times
- Yasmin's Theory: I don't trust this guy, I think heads will come up more often than tails


"Suit yourself -- here, take the coin and do whatever tests you want with it!"

What do the two theories predict in terms of the outcome of a series of coin flips?


```
In [81]: x_predictions = [0.5, 0.5]
         y_predictions = [0.4, 0.6]

         import matplotlib.pyplot as plt
         def plot_prediction(prediction, who):
             plt.bar([0,1], prediction)
             plt.ylim([0,1])
             plt.title(f"{who}'s coin flip prediction")
             plt.show()
```

```
In [87]: plot_prediction(x_predictions, "Xavier")
```



```
In [88]: plot_prediction(y_predictions, "Yasmin")
```



- Data:

```
In [83]: from collections import Counter
import pandas as pd
import numpy as np
import secret_coin

def do_coin_flips(N):
    coin_flips = np.array([secret_coin.flip() for i in range(N)])
    return coin_flips

def get_flip_distributions(x_predictions, y_predictions, results):
    flip_counter = Counter(results)
    flip_counts = np.array([flip_counter[0], flip_counter[1]]) / len(results)
    flip_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': flip_counts, 'which': ['Actual outcome', 'Actual outcome']})
    x_pred_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': x_predictions, 'which': ["Xavier's prediction", "Xavier's prediction"]})
    y_pred_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': y_predictions, 'which': ["Yusuf's prediction", "Yusuf's prediction"]})
    full_df = pd.concat([flip_df, x_pred_df, y_pred_df]).reset_index()
    return full_df
```



```
In [84]: flips_100 = do_coin_flips(100)
flips_100
```

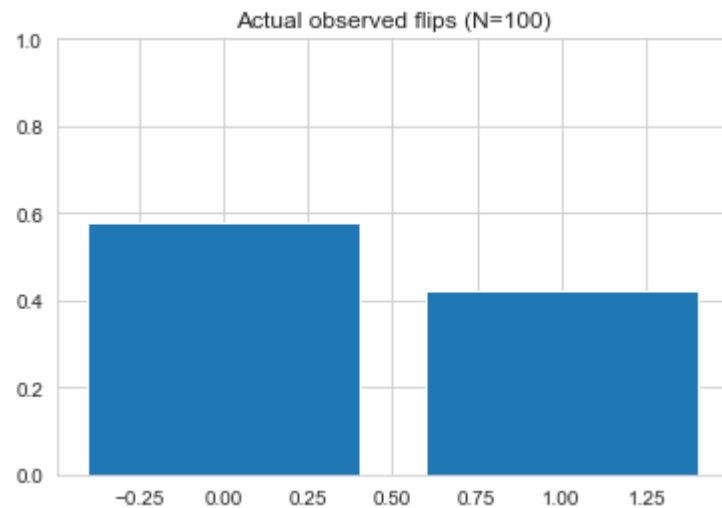
```
Out[84]: array([1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0,
        1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1,
        1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
        1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1,
        1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1])
```

```
In [85]: dist_df = get_flip_distributions(x_predictions, y_predictions, flips_100)
dist_df
```

```
Out[85]:
```


	index	outcome	p	which
0	0	Tails	0.58	Actual outcome
1	1	Heads	0.42	Actual outcome
2	0	Tails	0.50	Xavier's prediction
3	1	Heads	0.50	Xavier's prediction
4	0	Tails	0.40	Yusuf's prediction
5	1	Heads	0.60	Yusuf's prediction

```
In [86]: plt.bar([0,1], [dist_df.iloc[0]['p'],dist_df.iloc[1]['p']])  
plt.title("Actual observed flips (N=100)")  
plt.ylim([0,1])  
plt.show()
```




But... **how** wrong is our prediction?

```
In [90]: import seaborn as sns; sns.set_style("whitegrid")
def plot_distributions_a(dist_df):
    g = sns.catplot(data=dist_df, kind="bar", x="outcome", y="p", hue="which")
    g.despine(left=True)
    g.set_axis_labels("Outcome", "P(outcome)")
    g.legend.set_title("")
    g.set(ylim=(0,1))
    plot_distributions_a(dist_df)
```



```
In [70]: def plot_distributions_b(dist_df):  
    g = sns.catplot(  
        data=dist_df, kind="bar",  
        x="which", y="p", hue="outcome"  
    )  
    g.despine(left=True)  
    g.set_axis_labels("Which Distribution?", "P(outcome)")  
    g.legend.set_title("Outcome")  
    g.set(ylim=(0,1))  
plot_distributions_b(dist_df)
```



Hmm... that actual outcome still looks kinda sketchy. Let's try one more time

```
In [71]: flips_100_2 = do_coin_flips(100)
flips_100_2
```


```
Out[71]: array([0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
        1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0,
        1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1,
        0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0,
        1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0])
```

```
In [72]: dist_df2 = get_flip_distributions(x_predictions, y_predictions, flips_100_2)
dist_df2
```


Out[72]:

	index	outcome	p	which
0	0	Tails	0.51	Actual outcome
1	1	Heads	0.49	Actual outcome
2	0	Tails	0.50	Xavier's prediction
3	1	Heads	0.50	Xavier's prediction
4	0	Tails	0.40	Yusuf's prediction
5	1	Heads	0.60	Yusuf's prediction

```
In [73]: plot_distributions_a(dist_df2)
```



```
In [74]: plot_distributions_b(dist_df2)
```



So, we need to **formalize** how to measure **how bad** a prediction is.

Enter... statistics!

Null vs. Alternative Hypotheses

- **Null Hypothesis (H_0):** The skeptical hypothesis... "Nothing interesting is going on here, any patterns were simply due to chance"
 - The coin is not weird. $P(\text{heads}) = 0.5$
- **Alternative Hypothesis (H_A):** Something *other than* chance is generating the pattern we observe
 - The coin is loaded! $P(\text{heads}) \neq 0.5$

ONLY TWO POSSIBLE CONCLUSIONS FROM YOUR HYPOTHESIS TEST

1. "We **reject** the null hypothesis"
2. "We **fail to reject** the null hypothesis"

Test Statistic

- Computed from the **observed** data
- A measure of how reasonable our alternative hypothesis is for explaining this data

(This is the number we were looking for before!)

So, how bad were our coin flip predictions?

```
In [91]: def test_stat_a(coin_flips):  
         # Num heads - num tails  
         num_heads = len([f for f in coin_flips if f == 1])  
         num_tails = len([f for f in coin_flips if f == 0])  
         return num_heads - num_tails
```

```
In [92]: print(test_stat_a(do_coin_flips(100)))  
         print(test_stat_a(do_coin_flips(100)))
```

-2

4

What would this test statistic look like if the coin was *actually* fair?


```
In [93]: def fair_coin_flips(N):  
         return np.array([np.random.binomial(1, 0.5) for i in range(N)])
```

```
In [94]: fair_coin_flips(100)
```

```
Out[94]: array([0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1,  
                0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1,  
                1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0,  
                0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1,  
                0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1])
```

```
In [95]: test_stats = []  
         for trial_num in range(1000):  
             test_stats.append(test_stat_a(fair_coin_flips(100)))
```

```
In [96]: plt.hist(test_stats)  
         plt.show()
```




So... now we do 100 flips of our secret coin and see where it falls on this distribution!

```
In [114]: secret_coin_results = do_coin_flips(100)
```

```
In [115]: sc_stat = test_stat_a(secret_coin_results)
sc_stat
```

```
Out[115]: 14
```

```
In [116]: plt.hist(test_stats)
plt.vlines(sc_stat, 0, 250, color='red')
plt.show()
```




Note: This is **not** the only possible test statistic!

- We *make up* the test statistic that can best help us "detect" not-by-chance data

```
In [117]: def test_stat_b(coin_flips):  
    # The squared difference in predicted probabilities  
    num_heads = len([f for f in coin_flips if f == 1])  
    num_tails = len([f for f in coin_flips if f == 0])  
    return (num_heads - num_tails) ** 2
```

```
In [118]: test_stats_b = []  
for trial_num in range(1000):  
    test_stats_b.append(test_stat_b(fair_coin_flips(100)))
```

```
In [119]: plt.hist(test_stats_b)  
plt.show()
```




Now let's place our **observed** data on this plot

```
In [137]: sc_stat_b = test_stat_b(secret_coin_results)
sc_stat_b
```

```
Out[137]: 196
```


```
In [139]: plt.hist(test_stats_b)
plt.vlines(sc_stat_b, 0, 800, color='red')
plt.show()
```



```
In [141]: def test_stat_c(coin_flips):  
    # 1 if it's within 5, 0 otherwise  
    num_heads = len([f for f in coin_flips if f == 1])  
    num_tails = len([f for f in coin_flips if f == 0])  
    diff = abs(num_heads - num_tails)  
    return 1 if diff <= 5 else 0
```

```
In [128]: test_stats_c = []  
    for trial_num in range(1000):  
        test_stats_c.append(test_stat_c(fair_coin_flips(100)))
```


```
In [129]: plt.hist(test_stats_c)  
    plt.show()
```




```
In [130]: sc_stat_c = test_stat_c(secret_coin_results)
sc_stat_c
```

```
Out[130]: 0
```

```
In [136]: plt.hist(test_stats_c, align='left')
plt.vlines(sc_stat_c, 0, 800, color='red')
plt.show()
```



The Normal Distribution



Correlation

