

Week 8: Hypothesis Testing, Correlation vs. Causation

DSUA111: Data Science for Everyone, NYU, Fall 2020

TA Jeff, jpj251@nyu.edu

- This slideshow: <https://jjacobs.me/dsua111-sections/week-08>
(<https://jjacobs.me/dsua111-sections/week-08>).
- All materials: <https://github.com/jpowerj/dsua111-sections>
(<https://github.com/jpowerj/dsua111-sections>).

Outline

1. Hypothesis Testing Overview
2. Testing Coins
3. Null vs. Alternative Hypotheses
4. Test Statistics
5. The Normal Distribution
6. Correlation vs. Causation

Hypothesis Testing Overview

tl;dr

- If your theory was true, what would the data look like?
- Now compare that to the actual, **observed** data

Testing Coins

Example: I walk up to you and say "Hey, wanna gamble? We'll each put in a dollar, then I'll flip this coin. Heads I get the \$2, tails you get the \ \$2"

- Xavier's Theory: I think the coin is fair. Heads and tails will come up about the same number of times
- Yasmin's Theory: I don't trust this guy, I think heads will come up more often than tails

"Suit yourself -- here, take the coin and do whatever tests you want with it!"

What do the two theories predict in terms of the outcome of a series of coin flips?

```
In [ ]: x_predictions = [0.5, 0.5]
        y_predictions = [0.4, 0.6]

import matplotlib.pyplot as plt
def plot_prediction(prediction, who):
    plt.bar([0,1], prediction)
    plt.ylim([0,1])
    plt.title(f"{who}'s coin flip prediction")
    plt.show()
```

```
In [ ]: plot_prediction(x_predictions, "Xavier")
```

```
In [ ]: plot_prediction(y_predictions, "Yasmin")
```

- Data:

```
In [ ]: from collections import Counter
import pandas as pd
import numpy as np
import secret_coin
np.random.seed(1948)

def do_coin_flips(N):
    coin_flips = np.array([secret_coin.flip() for i in range(N)])
    return coin_flips

def get_flip_distributions(x_predictions, y_predictions, results):
    flip_counter = Counter(results)
    flip_counts = np.array([flip_counter[0], flip_counter[1]]) / len(results)
    flip_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': flip_counts, 'which': ['Actual outcome', 'Actual outcome']})
    x_pred_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': x_predictions, 'which': ["Xavier's prediction", "Xavier's prediction"]})
    y_pred_df = pd.DataFrame({'outcome': ["Tails", "Heads"], 'p': y_predictions, 'which': ["Yasmin's prediction", "Yasmin's prediction"]})
    full_df = pd.concat([flip_df, x_pred_df, y_pred_df]).reset_index()
    return full_df
```



```
In [ ]: flips_100 = do_coin_flips(100)
        flips_100
```

```
In [ ]: dist_df = get_flip_distributions(x_predictions, y_predictions, flips_100)
        dist_df
```

```
In [ ]: plt.bar([0,1], [dist_df.iloc[0]['p'],dist_df.iloc[1]['p']])  
plt.title("Actual observed flips (N=100)")  
plt.ylim([0,1])  
plt.show()
```

But... **how** wrong is our prediction?

```
In [ ]: import seaborn as sns; sns.set_style("whitegrid")
def plot_distributions_a(dist_df):
    g = sns.catplot(data=dist_df, kind="bar", x="outcome", y="p", hue="which")
    g.despine(left=True)
    g.set_axis_labels("Outcome", "P(outcome)")
    g.legend.set_title("")
    g.set(ylim=(0,1))
plot_distributions_a(dist_df)
```

```
In [ ]: def plot_distributions_b(dist_df):  
        g = sns.catplot(  
            data=dist_df, kind="bar",  
            x="which", y="p", hue="outcome"  
        )  
        g.despine(left=True)  
        g.set_axis_labels("Which Distribution?", "P(outcome)")  
        g.legend.set_title("Outcome")  
        g.set(ylim=(0,1))  
plot_distributions_b(dist_df)
```

Hmm... that actual outcome still looks kinda sketchy. Let's try one more time

```
In [ ]: flips_100_2 = do_coin_flips(100)
        flips_100_2
```

```
In [ ]: dist_df2 = get_flip_distributions(x_predictions, y_predictions, flips_100_2)
        dist_df2
```

```
In [ ]: plot_distributions_a(dist_df2)
```

```
In [ ]: plot_distributions_b(dist_df2)
```

So, we need to **formalize** how to measure **how bad** a prediction is.

Enter... statistics!

Null vs. Alternative Hypotheses

- **Null Hypothesis (H_0):** The skeptical hypothesis... "Nothing interesting is going on here, any patterns were simply due to chance"
 - The coin is not weird. $P(\text{heads}) = 0.5$
- **Alternative Hypothesis (H_A):** Something *other than* chance is generating the pattern we observe
 - The coin is loaded! $P(\text{heads}) \neq 0.5$

ONLY TWO POSSIBLE CONCLUSIONS FROM YOUR HYPOTHESIS TEST

1. "We **reject** the null hypothesis"
2. "We **fail to reject** the null hypothesis"

Test Statistic

- Computed from the **observed** data
- A measure of how reasonable our alternative hypothesis is for explaining this data

(This is the number we were looking for before!)

So, how bad were our coin flip predictions?

```
In [ ]: def test_stat_a(coin_flips):  
        # Num heads - num tails  
        num_heads = len([f for f in coin_flips if f == 1])  
        num_tails = len([f for f in coin_flips if f == 0])  
        return num_heads - num_tails
```

```
In [ ]: print(test_stat_a(do_coin_flips(100)))  
        print(test_stat_a(do_coin_flips(100)))
```

What would this test statistic look like if the coin was *actually* fair?

```
In [ ]: def fair_coin_flips(N):  
        return np.array([np.random.binomial(1, 0.5) for i in range(N)])
```

```
In [ ]: fair_coin_flips(100)
```

```
In [ ]: test_stats = []  
        for trial_num in range(1000):  
            test_stats.append(test_stat_a(fair_coin_flips(100)))
```

```
In [ ]: plt.hist(test_stats)  
        plt.show()
```

So... now we do 100 flips of our secret coin and see where it falls on this distribution!

```
In [ ]: secret_coin_results = do_coin_flips(100)
```

```
In [ ]: sc_stat = test_stat_a(secret_coin_results)
sc_stat
```

```
In [ ]: plt.hist(test_stats)
plt.vlines(sc_stat, 0, 250, color='red')
plt.show()
```

Note: This is **not** the only possible test statistic!

- We *make up* the test statistic that can best help us "detect" not-by-chance data

```
In [ ]: def test_stat_b(coin_flips):  
        # The squared difference in predicted probabilities  
        num_heads = len([f for f in coin_flips if f == 1])  
        num_tails = len([f for f in coin_flips if f == 0])  
        return (num_heads - num_tails) ** 2
```

```
In [ ]: test_stats_b = []  
        for trial_num in range(1000):  
            test_stats_b.append(test_stat_b(fair_coin_flips(100)))
```

```
In [ ]: plt.hist(test_stats_b)  
        plt.show()
```


Now let's place our **observed** data on this plot

```
In [ ]: sc_stat_b = test_stat_b(secret_coin_results)
        sc_stat_b
```

```
In [ ]: plt.hist(test_stats_b)
        plt.vlines(sc_stat_b, 0, 800, color='red')
        plt.show()
```

```
In [ ]: def test_stat_c(coin_flips):  
        # 1 if it's within 5, 0 otherwise  
        num_heads = len([f for f in coin_flips if f == 1])  
        num_tails = len([f for f in coin_flips if f == 0])  
        diff = abs(num_heads - num_tails)  
        return 1 if diff <= 5 else 0
```


```
In [ ]: test_stats_c = []  
        for trial_num in range(1000):  
            test_stats_c.append(test_stat_c(fair_coin_flips(100)))
```

```
In [ ]: plt.hist(test_stats_c)  
        plt.show()
```

```
In [ ]: sc_stat_c = test_stat_c(secret_coin_results)
        sc_stat_c
```

```
In [ ]: plt.hist(test_stats_c, align='left')
        plt.vlines(sc_stat_c, 0, 800, color='red')
        plt.show()
```

The Normal Distribution



Correlation

